

EXPERTMENT - 5

MAPREDUCE PROGRAM TO FIND THE GRADES OF STUDENTS

*AIM: To develop a mapreduce program to find the grades of students

*PROGRAM:

```
import java.util.Scanner;
public class JavaExample
{
    public static void main (String args[])
    {
        int marks[] = new int[6];
        int i;
        float total = 0, avg;
        Scanner scanner = new Scanner (System.in);
        for (i=0; i<6; i++)
        {
            System.out.print ("Enter Marks of subject " +
                (i+1) + ": ");
            marks[i] = scanner.nextInt();
            total = total + marks[i];
        }
        scanner.close();
        // here avg = total/6;
        System.out.print ("The student Grade is: ");
    }
}
```

```
if (avg >= 80)
{
    System.out.print ("A");
}
else if (avg >= 60 && avg < 80)
{
    System.out.print ("B");
}
else if (avg >= 40 && avg < 60)
{
    System.out.print ("C");
}
else
{
    System.out.print ("D");
}
```

OUTPUT:

Enter marks of Subject 1 : 40

Enter marks of Subject 2 : 80

Enter marks of Subject 3 : 80

Enter marks of Subject 4 : 40

Enter marks of Subject 5 : 60

Enter marks of Subject 6 : 60

The student Grade is B

EXPERIMENT - 6

*AIM: Retrieving user login credentials from /etc/passwd using

*PROGRAM:

⇒ Step 1: Copy /etc/passwd into pig-programs folder as passwd.txt

⇒ Step 2: Create pass-analysis.pig

```
pass = load "passwd.txt" using pigStorage("::") as  
(username:chararray, passwd:chararray, userid:int,  
groupid:int, userinfo:chararray, root_dir:chararray  
shell:chararray);
```

```
user_list = foreach pass generate username, passwd;  
dump user_list;
```

⇒ Step 3: Create pass-analysis1.pig

```
pass = load "passwd.txt" using pigStorage("::") as  
(username:chararray, passwd:chararray, userid:int  
root_dir:chararray);
```

```
grp_by_shell = group pass by shell;
```

```
STORE pass INTO 'pig_output' USING pigStorage('::');
```

Step 4: Pig -> local pars.analysis.pig

Step 5: Pig -> local pars.analysis1.pig

Step 6: Result can be seen through command interface as
part-r 00000 \$hadoop dfs -cat Pig-output/

EXPERIMENT - 3

* AIM: Working with hive

* HIVE PROGRAMS:

⇒ \$ Hive:

↳ Create Database

Syntax: CREATE DATABASE | SCHEMA [IF NOT EXISTS]
<database name>

hive > CREATE DATABASE [IF NOT EXISTS] userdb;

or

hive > CREATE SCHEMA userdb;

↳ Drop Database

Syntax: DROP DATABASE statement DROP [DATABASE |
SCHEMA] [IF EXISTS] database_name
[RESTRICT | CASCADE];

hive > DROP DATABASE IF EXISTS userdb;

hive > DROP DATABASE IF EXISTS userdb CASCADE;

hive > DROP SCHEMA userdb;

↳ Create Table

Syntax: CREATE [TEMPORARY] [EXTERNAL] TABLE [IF NOT
EXISTS] [db_name.]

[Col_name data_type [comment col_comment], ...]

[COMMENT table comment]

[ROW FORMAT row format]

[STORED AS file-format]

hive > CREATE TABLE IF NOT EXISTS employee (eid int,
name String, salary Double, designation String,
department String)
COMMENT 'Employee details'
ROW FORMAT DELIMITED
STORED AS TEXTFILE;

↳ Create table from existing table

Create table emp select * from employee;

↳ Creating table with buckets:

CREATE TABLE students (sid int, sname String, age
INT, dept String, gpa Decimal(3,2))
CLUSTERED BY (dept) INTO 4 BUCKETS
STORED AS ORC;

↳ ALTER TABLE:

Syntax =

ALTER TABLE name RENAME TO new-name

ALTER TABLE name ADD COLUMNS (col-spec [, col-spec])

ALTER TABLE name DROP [COLUMN] column_name

Rename To.

hive > ALTER TABLE employee RENAME

Change

hive > ALTER TABLE employee CHANGE name ename
String;

Add Column:

hive > ALTER TABLE employee ADD COLUMNS

Replace:

hive ALTER TABLE employee REPLACE COLUMNS
(empid INT, ename STRING)

↳ DROP TABLE:

Syntax

DROP TABLE [IF EXISTS] table_name;

hive > DROP TABLE IF EXISTS employee

↳ loading files into Tables:

Syntax:

LOAD DATA [LOCAL] INPATH "filepath" [OVERWRITE] INTO
TABLE tablename [PARTITION (partcol1=val1, partcol2=
val2)]

Insert Data from a table:

INSERT OVERWRITE TABLE tablename [PARTITION (partcol1=
val1, partcol2=val2 ...)] [IF NOT EXISTS]

* HIVE - Partitioning:

Date :

Hive partitions is a way to organize tables into partitions by dividing tables into different parts.

1. Creation of Table studentInfo

Create table studentInfo (sid int, sname String,
phoneno BIGINT, dept String)

2. Loading data into created table all states

3. Creation of partition table.

4. For partition we have to set this property set hive.exec.dynamic.partition=true

INSERT OVERWRITE TABLE student_part PARTITION(dept)
SELECT * from studentInfo;

a) Adding a partition using Alter Table

Syntax:

ALTER TABLE table_name ADD [IF NOT EXISTS] PARTITION
partition_spec [LOCATION 'location'] partition_spec
partition_spec (p_column = p_col_value, p_column = p_col_value)
hive> ALTER TABLE student_part ADD PARTITION

b. DROPPING A PARTITION:

Syntax:

ALTER TABLE table_name DROP [IF EXISTS] PARTITION
partition_spec, PARTITION partition_spec...;

Merge two tables

Standard Syntax:

```
MERGE INTO <target table> AS T USING <Source expression
/> AS S
```

```
ON <boolean expression1>
```

```
WHEN MATCHED [AND <boolean expression2>] THEN
```

```
UPDATE SET <set clause list>
```

```
WHEN MATCHED [AND <boolean expression3>] THEN DELETE
```

* HiveQL

Syntax:

```
SELECT [ALL | DISTINCT] select-epr, select-expr...
```

```
FROM table reference
```

```
[WHERE where condition]
```

```
[GROUP BY col-list]
```

```
[HAVING having condition]
```

```
[ORDER BY col-list]]
```

```
[Limit number];
```

1. Query to retrieve the employee details whose Id is 1205

```
hive > SELECT * FROM employee WHERE Id = 1205
```

2. Query to retrieve the employee details whose salary is more than or equal Rs 4000

```
hive > SELECT * FROM employee WHERE salary >= 4000;
```

* Simple Queries on employee table.

Id	Name	Salary	Designation	Dept
1201	Gopal	10000	Technical manager	TP
1202	Manisha	10000	Marketing	PR
1203	Kiran	10000	Hr Admin	HR
1204	Kranthi	10000	Op Admin	Admin

3. Query to retrieve employee details whose dept is TP and salary is more than RS 4000

hive > SELECT * FROM employee WHERE salary > 4000
and dept = TP;

* Aggregate Functions :

hive > SELECT count(*) FROM employee;

hive > SELECT sum(Salary) FROM employee;

hive > SELECT avg(Salary) FROM employee;

* Hive - Built - in functions :

Create a dummy table with single row and column &
insert a value into it.

Hive > Create table dummy (value String);

Hive > load data local inpath '/home/harunin/gdummydata' into table dummy;

1) round() function

hive > SELECT round(2.6) from dummy;

2) floor() function

hive > SELECT floor(2.6) from dummy;

ADD	
50	

* VIEW AND INDEXES

→ Creating a View:

↳ Syntax:

CREATE VIEW [IF NOT EXISTS] view_name [(column_name
[COMMENT column_comment], ...)]

Query to create a view named as emp_30000 with
Salary > 30000

hive> CREATE VIEW emp_30000 AS SELECT * FROM employee
WHERE salary > 30000;

→ Dropping a View:

↳ Syntax:

DROP VIEW view_name

Query to drops a view named as emp_30000;

hive> DROP VIEW emp_30000;

Creating an Index:

↳ Syntax:

CREATE INDEX index_name

ON TABLE base_table_name (col_name, ...)

AS 'index_handler.class.name'

[WITH DEFERRED REBUILD]

[IN TABLE index_table_name]

[

[ROW FORMAT...] STORED AS ...

]

Testing:

hive > SELECT avg(Salary) FROM employee;

hive > SHOW FORMATTED INDEX ON employee;

hive > CREATE INDEX in_bit_empsalary ON TABLE employee(Salary)

hive > SELECT avg(Salary) FROM employee;

AS "org.apache.hadoop.hive.q1.index.bitmap.BitmapIndexHandler";

hive > ALTER INDEX in_bit_empsalary ON employee
REBUILD;

hive > SELECT avg(Salary) FROM employee;

Syntax:

DROP INDEX <index_name> ON <table_name>

The following query drops an index named index_salary.

* JOINS:

↳ Syntax:

join_table:

table_reference JOIN table_factor [join_condition]

| table_reference {LEFT|RIGHT|FULL} [C OUTER] JOIN

table reference

Join condition

| table_reference LEFT SEMI JOIN table_reference join
condition

JOIN on the CUSTOMER and ORDER tables

hive > SELECT c.ID, c.NAME, c.AGE, o.AMOUNT
FROM CUSTOMERS c JOIN ORDERS ON (c.ID=o.CUSTOMER_ID)
LEFT OUTER JOIN between CUSTOMER and ORDER tables

hive > SELECT c.ID, c.NAME, o.AMOUNT
FROM CUSTOMERS c FULL OUTER JOIN ORDERS ON
c.ID = o.CUSTOMER_ID;

Example

Consider the following table named CUSTOMERS

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000.00
2	Khilan	25	Delhi	1500.00
3	Hanika	27	Kota	2000.00
4	Muffy	24	Mumbai	6500.00

⇒ Consider another table ORDERS as follows:

OID	customer-ID	Amount
102	3	3000
100	3	1500
101	2	1500

EXPERIMENT - 8

* AIM: Hive UDF (User Defined functions)

HIVE UDF PROGRAMS:

Steps to write a UDF function in Hive

1. Create a java class for the User defined function which extends org.apache.hadoop.hive.sql.exec.UDF and implements more than one evaluate()
2. Package your java class into a JAR file
3. Go to Hive CLI, add your JAR and verify your JAR's
4. Use it in Hive SQL

↳ Simple UDF for "trim-like" function called "strip"

1. Open Eclipse and create a new project
2. Create a new package with name shipword
3. Create a new class file with name strip
 - ↳ Hadoop/Hadoop-core.jar
 - ↳ Hadoop/lib/Common - cli - 1.2.jar

PROGRAM:

Package shipword;

```
import org.apache.commons.lang.StringUtil;
```

```
import org.apache.hadoop.hive.udf.exec.UDF;  
import org.apache.hadoop.io.Text;
```

```
public class Strip extends UDF {
```

```
    private Text result = new Text();
```

```
    public Text evaluate(Text str, String stripChars) {
```

```
        if (str == null) {
```

```
            return null;
```

```
}
```

```
        result.set(StringUtil.strip(str.toString(), stripChars));
```

```
        return result;
```

```
}
```

```
public Text evaluate(Text str) {
```

```
    if (str == null) {
```

```
        return null;
```

```
}
```

```
    result.set(StringUtil.strip(str.toString()));
```

```
    return result;
```

```
4
```

```
3
```

Verify JAR is in Hive classpath

```
hive> list jars:
```

```
/usr/lib/hive/lib/hive-contrib.jar
```

```
/home/training/HiveUDFs/stripingword.jar
```

2. Convert celcius to fahrenheit

Package udfpack;

```
import org.apache.hadoop.hive.udf.UDF;
```

```
public class ConvertToCelcius extends UDF {
    public double evaluate (double value) {
        return (value - 32) / 1.8;
    }
}
```

3

```
hive> add jar /home/training/HiveUDFs/convert.jar
hive> SELECT ConvertToCelcius(10) from dummy;
```

2. Overloading "evaluate" method

Package udfpack;

```
import org.apache.hadoop.hive.udf.UDF;
```

```
public class AbsValue extends UDF {
    public Double evaluate (double value) {
        return Math.abs(value);
    }
}
```

3

```
public long evaluate (long value) {
    return Math.abs(value);
```

4

```
hive> addjar /home/training/HiveUDFs/absval.jar
```

hive > SELECT AbsValue(20.65) from dummy;

3. Say hello:

Package udfpack:

```
import org.apache.hadoop.hive.ql.exec.UDF;  
import org.apache.hadoop.io.Text;
```

class SayHello extends UDF

{

public Text evaluate(Text Input)

{

return new Text("Hello" + input.toString());

}

}

hive > add jar / home/training/HiveUDFs/sayHello.jar

hive > create temporary function sayHello as 'udfpack.SayHello';

hive > SELECT sayHello('Hadoop') from dummy;

4. Convert to lower case:

```
import org.apache.hadoop.hive.ql.exec.UDF;
```

```
import org.apache.hadoop.io.Text;
```

public final class Lower extends UDF {

```
public Text evaluate(final Text s){  
    if(s==null) {return null;}  
    return new Text(s.toString().toLowerCase());
```

{

}

```
hive> add jar /home/training/HiveUDFs/lower.jar  
hive> SELECT lower('HADOOP') from dummy;  
hive > SELECT lower(name) from employee;
```

UDF OUTPUT:

hive> select strip('hadoop','ha') from dummy;

OK

doop

Time taken: 0.131 seconds, Fetched: 1 row(s)

hive> select strip

OK

hive UDF